

MLCC Laboratory 3: Dimensionality reduction, feature selection

In this laboratory we will address the problem of data analysis and dimensionality reduction with a reference to a classification problem.

Think hard before you call the instructors or you look at the solution file!

1 Warm up - data generation

You will generate a training and a test set of d -dimensional points (n points for each class), with $n = 100$ and $d = 30$. Only two of those dimensions will be meaningful, the other one will be a variable we will modify.

- **1.A** For each point, the first two variables will be generated by `MixGauss`, extracted from two gaussian distributions with centroids $(1, 1)$ and $(-1, -1)$ and standard deviation 0.7 (the first one with labels 1 , the second with label -1):

```
1 Xtr, Ytr = MixGauss(means=[[0, 0], [1, 1]], sigmas=[[0.7], [0.7]], n=100)
2 Ytr = 2*np.mod(Ytr, 2)-1
3 Xts, Yts = MixGauss(means=[[0, 0], [1, 1]], sigmas=[[0.7], [0.7]], n=100)
4 Yts = 2*np.mod(Yts, 2)-1
```

- **1.B.** You may want to plot the relevant variables of the data:

```
1 fig, axs = plt.subplots(1, 1)
2 plt.scatter(Xtr[:, 0], Xtr[:, 1], s=30, c=np.squeeze(Ytr), alpha=0.8)
3 plt.scatter(Xts[:, 0], Xts[:, 1], s=30, c=np.squeeze(Ytr), alpha=0.1)
4 plt.title('train and test datasets')
```

- **1.C** The remaining variables will be generated as gaussian noise:

```
1 sigma_noise = 0.01
2 Xtr_noise = sigma_noise * np.random.randn(2*n, d-2)
3 Xts_noise = sigma_noise * np.random.randn(2*n, d-2)
```

To compose the final data matrix, run:

```
1 Xtr = np.concatenate((Xtr, Xtr_noise), axis=1)
2 Xts = np.concatenate((Xts, Xts_noise), axis=1)
```

2 Principal Component Analysis (PCA)

- **2.A** Compute the data principal components (see "`help(PCA)`").
- **2.B** Plot the first two components of `X_proj` using the following line:

```

1 fig, axs = plt.subplots(1, 1)
2 plt.scatter(X_proj[:, 0], X_proj[:, 1], s=30, c=np.squeeze(Ytr), alpha
    =0.8)
3 plt.title('train dataset projected on first 2 components')

```

- **2.C** Plot the cumulative sum of the eigenvalues you found.

Reason on the meaning of the results you are obtaining.

- **2.D** Display the `sqrt` of the first 10 eigenvalues and plot the coefficients (eigenvector) associated with the largest eigenvalue:

```

1 print(np.sqrt(D[:10]))
2 fig, axs = plt.subplots(1, 1)
3 plt.scatter(range(d), V[:, 0], s=30, alpha=0.8)
4 plt.title('Eigenvector of highest eigenvalue')

```

- **2.E** Repeat the above steps with dataset generated using different `sigma_noise = 0, 0.01, 0.1, 0.5, 0.7, 1, 1.2, 1.4, 1.6, 2`. To what extent data visualization by PCA is affected by the noise?

3 Variable selection

- **3.A** Use the data generated in part 1. Standardize the data matrix, so that each column has mean 0 and standard deviation 1:

```

1 m = np.mean(Xtr, axis=0)
2 s = np.std(Xtr, axis=0)
3 Xtr = (Xtr - m) / s

```

Do the same for `Xts`, by using `m` and `s` computed on `Xtr`.

- **3.B** Use the orthogonal matching pursuit algorithm (type `"help(OMatchingPursuit)"`).
- **3.C** You may want to check the predicted labels on the training set:

```

1 Ypred = np.sign(Xts.dot(w))
2 error = calcErr(Yts, Ypred)

```

and plot the coefficients `w` with `scatter(range(d), abs(w))`. How the error changes with the number of iterations of the method?

- **3.D** By using the method `holdoutCVOMP` find the best number of iterations with `intIter = 2, ..., d` (and, for instance, `perc = 0.75, nrip = 20`).

Moreover, plot the training and validation error with the following lines:

```

1 fig, axs = plt.subplots(1, 1)
2 plt.plot(intIter, Tm)
3 plt.plot(intIter, Vm)
4 plt.legend(['Training error', 'Validation error'])
5 plt.xlabel('number of iterations for OMP')
6 plt.ylabel('error')

```

What is the behavior of the training and the validation errors with respect to the number of iterations?

- **3.E** Try to increase the number of relevant variables $d = 3, 5, \dots$ (and the corresponding standard deviation of the Gaussians) around the centroids:

```

1 np.ones((d, 1)) # vector of all 1s
2 # and
3 -np.ones((d, 1)) # vector of all -1s

```

and see how this change is reflected in the cross-validation.

4 If you have time - more experiments

- **4.A** Analyse the results you obtain on sections 2 and 3 once you choose:

- $n \gg d$
- $n \approx d$
- $n \ll d$,

and evaluate the benefits of the two different analyses.

- **4.B** Dimensionality reduction is often used as a pre-processing step to a learning (classification) algorithm. The idea is to perform classification in a lower dimensional space and therefore save computational time. You have the following task:
 - Generate a new training and test datasets as in Laboratory 1.
 - Perform dimensionality reduction on the training set.
 - Using the projection you just found, project the test set.
 - Perform kNN in the lower dimensional (projected) space. Compare the result (both accuracy and running time) with the one in Laboratory 1.